

Question 11.1

Using the crime data set `uscrime.txt` from Questions 8.2, 9.1, and 10.1, build a regression model using:

1. Stepwise regression
2. Lasso
3. Elastic net

For Parts 2 and 3, remember to scale the data first – otherwise, the regression coefficients will be on different scales and the constraint won't have the desired effect.

For Parts 2 and 3, use the `glmnet` function in R.

Notes on R:

- For the elastic net model, what we called λ in the videos, `glmnet` calls “alpha”; you can get a range of results by varying alpha from 1 (lasso) to 0 (ridge regression) [and, of course, other values of alpha in between].
- In a function call like `glmnet(x,y,family="mgaussian",alpha=1)` the predictors `x` need to be in R's matrix format, rather than data frame format. You can convert a data frame to a matrix using `as.matrix` – for example, `x <- as.matrix(data[,1:n-1])`
- Rather than specifying a value of `T`, `glmnet` returns models for a variety of values of `T`.

Method 1. Stepwise Regression

The stepwise regression uses the Akaike Information Criterion (AIC) to add or remove variables iteratively to find the best-fitting model.

Code:

```
uscrime <- read.table("E:/Study/ISyE 6501/hw7/uscrime.txt", header = TRUE)

initial_model <- lm(Crime ~ ., data = uscrime)

sink("E:/Study/ISyE 6501/hw7/11.1.1_Output.txt")
stepwise_model <- step(initial_model, direction = "both")
print(stepwise_model)
summary(stepwise_model)
sink()

for (i in 1:6) {
  filename <- paste0("E:/Study/ISyE 6501/hw7/plot", i, ".png")
  png(filename)
  plot(stepwise_model, which = i, main = paste("Stepwise Model Plot - Plot", i))
  dev.off()
}
```

Iterations:

Start: AIC=514.65

Crime ~ M + So + Ed + Po1 + Po2 + LF + M.F + Pop + NW + U1 +
U2 + Wealth + Ineq + Prob + Time

	Df	Sum of Sq	RSS	AIC
- So	1	29 1354974	512.65	
- LF	1	8917 1363862	512.96	
- Time	1	10304 1365250	513.00	
- Pop	1	14122 1369068	513.14	
- NW	1	18395 1373341	513.28	
- M.F	1	31967 1386913	513.74	
- Wealth	1	37613 1392558	513.94	
- Po2	1	37919 1392865	513.95	
<none>		1354946	514.65	
- U1	1	83722 1438668	515.47	
- Po1	1	144306 1499252	517.41	
- U2	1	181536 1536482	518.56	
- M	1	193770 1548716	518.93	
- Prob	1	199538 1554484	519.11	
- Ed	1	402117 1757063	524.86	
- Ineq	1	423031 1777977	525.42	

Step: AIC=512.65

Crime ~ M + Ed + Po1 + Po2 + LF + M.F + Pop + NW + U1 + U2 +
Wealth + Ineq + Prob + Time

	Df	Sum of Sq	RSS	AIC
- Time	1	10341	1365315	511.01
- LF	1	10878	1365852	511.03
- Pop	1	14127	1369101	511.14
- NW	1	21626	1376600	511.39
- M.F	1	32449	1387423	511.76
- Po2	1	37954	1392929	511.95
- Wealth	1	39223	1394197	511.99
<none>			1354974	512.65
- U1	1	96420	1451395	513.88
+ So	1	29	1354946	514.65
- Po1	1	144302	1499277	515.41
- U2	1	189859	1544834	516.81
- M	1	195084	1550059	516.97
- Prob	1	204463	1559437	517.26
- Ed	1	403140	1758114	522.89
- Ineq	1	488834	1843808	525.13

Step: AIC=511.01

Crime ~ M + Ed + Po1 + Po2 + LF + M.F + Pop + NW + U1 + U2 +
Wealth + Ineq + Prob

	Df	Sum of Sq	RSS	AIC
- LF	1	10533	1375848	509.37
- NW	1	15482	1380797	509.54
- Pop	1	21846	1387161	509.75
- Po2	1	28932	1394247	509.99
- Wealth	1	36070	1401385	510.23
- M.F	1	41784	1407099	510.42
<none>			1365315	511.01
- U1	1	91420	1456735	512.05
+ Time	1	10341	1354974	512.65
+ So	1	65	1365250	513.00
- Po1	1	134137	1499452	513.41
- U2	1	184143	1549458	514.95
- M	1	186110	1551425	515.01
- Prob	1	237493	1602808	516.54
- Ed	1	409448	1774763	521.33
- Ineq	1	502909	1868224	523.75

Step: AIC=509.37

Crime ~ M + Ed + Po1 + Po2 + M.F + Pop + NW + U1 + U2 + Wealth +
Ineq + Prob

	Df	Sum of Sq	RSS	AIC
- NW	1	11675	1387523	507.77
- Po2	1	21418	1397266	508.09
- Pop	1	27803	1403651	508.31
- M.F	1	31252	1407100	508.42
- Wealth	1	35035	1410883	508.55

<none>			1375848	509.37
- U1	1	80954	1456802	510.06
+ LF	1	10533	1365315	511.01
+ Time	1	9996	1365852	511.03
+ So	1	3046	1372802	511.26
- Po1	1	123896	1499744	511.42
- U2	1	190746	1566594	513.47
- M	1	217716	1593564	514.27
- Prob	1	226971	1602819	514.54
- Ed	1	413254	1789103	519.71
- Ineq	1	500944	1876792	521.96

Step: AIC=507.77

Crime ~ M + Ed + Po1 + Po2 + M.F + Pop + U1 + U2 + Wealth + Ineq + Prob

	Df	Sum of Sq	RSS	AIC
- Po2	1	16706	1404229	506.33
- Pop	1	25793	1413315	506.63
- M.F	1	26785	1414308	506.66
- Wealth	1	31551	1419073	506.82
<none>			1387523	507.77
- U1	1	83881	1471404	508.52
+ NW	1	11675	1375848	509.37
+ So	1	7207	1380316	509.52
+ LF	1	6726	1380797	509.54
+ Time	1	4534	1382989	509.61
- Po1	1	118348	1505871	509.61
- U2	1	201453	1588976	512.14
- Prob	1	216760	1604282	512.59
- M	1	309214	1696737	515.22
- Ed	1	402754	1790276	517.74
- Ineq	1	589736	1977259	522.41

Step: AIC=506.33

Crime ~ M + Ed + Po1 + M.F + Pop + U1 + U2 + Wealth + Ineq + Prob

	Df	Sum of Sq	RSS	AIC
- Pop	1	22345	1426575	505.07
- Wealth	1	32142	1436371	505.39
- M.F	1	36808	1441037	505.54
<none>			1404229	506.33
- U1	1	86373	1490602	507.13
+ Po2	1	16706	1387523	507.77
+ NW	1	6963	1397266	508.09
+ So	1	3807	1400422	508.20
+ LF	1	1986	1402243	508.26
+ Time	1	575	1403654	508.31
- U2	1	205814	1610043	510.76

```
- Prob 1 218607 1622836 511.13
- M 1 307001 1711230 513.62
- Ed 1 389502 1793731 515.83
- Ineq 1 608627 2012856 521.25
- Po1 1 1050202 2454432 530.57
```

Step: AIC=505.07

Crime ~ M + Ed + Po1 + M.F + U1 + U2 + Wealth + Ineq + Prob

	Df	Sum of Sq	RSS	AIC
- Wealth	1	26493	1453068	503.93
<none>			1426575	505.07
- M.F	1	84491	1511065	505.77
- U1	1	99463	1526037	506.24
+ Pop	1	22345	1404229	506.33
+ Po2	1	13259	1413315	506.63
+ NW	1	5927	1420648	506.87
+ So	1	5724	1420851	506.88
+ LF	1	5176	1421398	506.90
+ Time	1	3913	1422661	506.94
- Prob	1	198571	1625145	509.20
- U2	1	208880	1635455	509.49
- M	1	320926	1747501	512.61
- Ed	1	386773	1813348	514.35
- Ineq	1	594779	2021354	519.45
- Po1	1	1127277	2553852	530.44

Step: AIC=503.93

Crime ~ M + Ed + Po1 + M.F + U1 + U2 + Ineq + Prob

	Df	Sum of Sq	RSS	AIC
<none>			1453068	503.93
+ Wealth	1	26493	1426575	505.07
- M.F	1	103159	1556227	505.16
+ Pop	1	16697	1436371	505.39
+ Po2	1	14148	1438919	505.47
+ So	1	9329	1443739	505.63
+ LF	1	4374	1448694	505.79
+ NW	1	3799	1449269	505.81
+ Time	1	2293	1450775	505.86
- U1	1	127044	1580112	505.87
- Prob	1	247978	1701046	509.34
- U2	1	255443	1708511	509.55
- M	1	296790	1749858	510.67
- Ed	1	445788	1898855	514.51
- Ineq	1	738244	2191312	521.24
- Po1	1	1672038	3125105	537.93

Call:

lm(formula = Crime ~ M + Ed + Po1 + M.F + U1 + U2 + Ineq + Prob,

```
data = uscrime)
```

Coefficients:

(Intercept)	M	Ed	Po1	M.F	U1	U2	Ineq
-6426.10	93.32	180.12	102.65	22.34	-6086.63	187.35	61.33
Prob							
-3796.03							

Call:

```
lm(formula = Crime ~ M + Ed + Po1 + M.F + U1 + U2 + Ineq + Prob,  
    data = uscrime)
```

Residuals:

Min	1Q	Median	3Q	Max
-444.70	-111.07	3.03	122.15	483.30

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	-6426.10	1194.61	-5.379	4.04e-06 ***
M	93.32	33.50	2.786	0.00828 **
Ed	180.12	52.75	3.414	0.00153 **
Po1	102.65	15.52	6.613	8.26e-08 ***
M.F	22.34	13.60	1.642	0.10874
U1	-6086.63	3339.27	-1.823	0.07622 .
U2	187.35	72.48	2.585	0.01371 *
Ineq	61.33	13.96	4.394	8.63e-05 ***
Prob	-3796.03	1490.65	-2.547	0.01505 *

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 195.5 on 38 degrees of freedom

Multiple R-squared: 0.7888, Adjusted R-squared: 0.7444

F-statistic: 17.74 on 8 and 38 DF, p-value: 1.159e-10

The initial model includes all variables: $\text{Crime} \sim \text{M} + \text{So} + \text{Ed} + \text{Po1} + \text{Po2} + \text{LF} + \text{M.F} + \text{Pop} + \text{NW} + \text{U1} + \text{U2} + \text{Wealth} + \text{Ineq} + \text{Prob} + \text{Time}$, and the starting AIC is 514.65. The stepwise procedure removes variables So, LF, Time, Pop, NW, M.F, Wealth, and Po2 one by one to see if AIC decreases.

At each step, the model tests whether removing a variable results in a lower AIC value. The variable is excluded if AIC decreases significantly after removing a variable (e.g., AIC reduces from 514.65 to 512.65 by removing So). The process continues until no further removal or addition of variables reduces AIC.

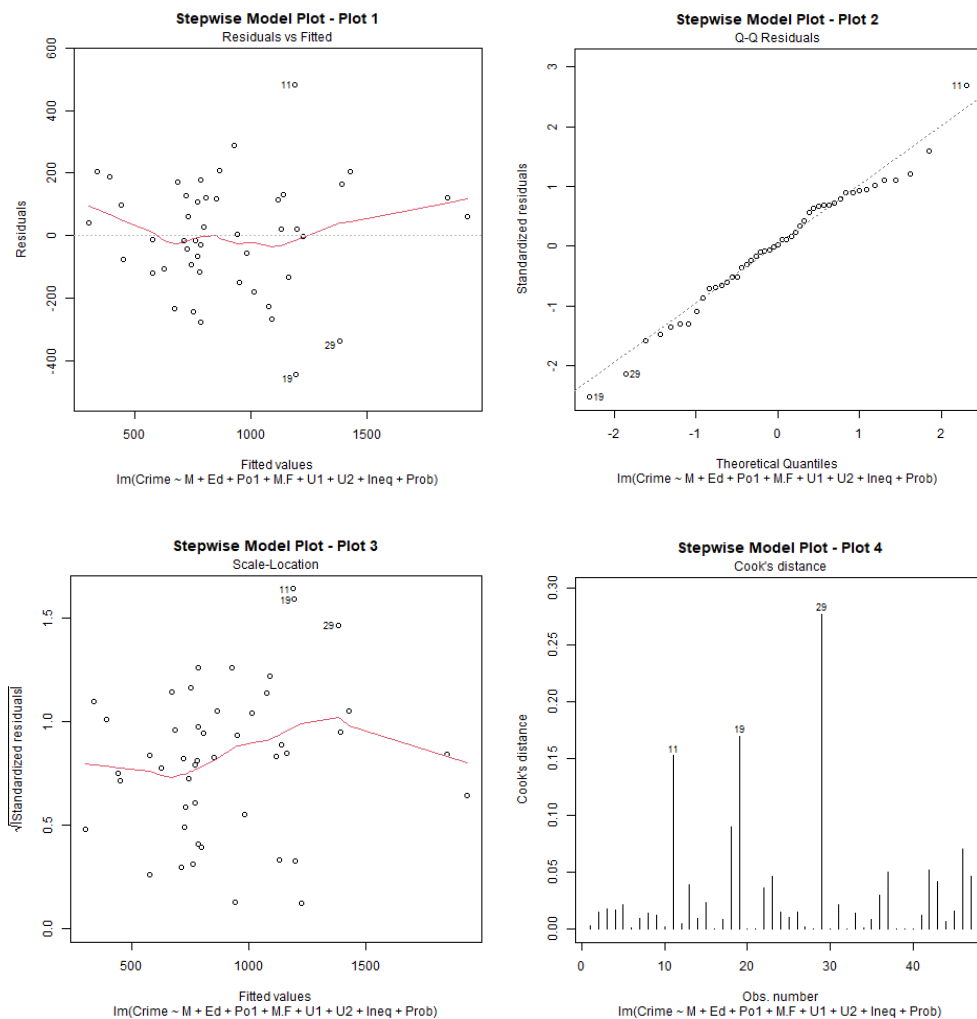
The final model includes $\text{Crime} \sim \text{M} + \text{Ed} + \text{Po1} + \text{M.F} + \text{U1} + \text{U2} + \text{Ineq} + \text{Prob}$. The final AIC value is 503.93, which is much lower than the initial AIC (514.65), indicating an improved model fit.


Thus, coefficients represent the estimated impact of each variable on the Crime rate. Significant variables include M, Ed, Po1, and Ineq, as shown by their low p-values (< 0.05), suggesting these variables strongly influence

the Crime rate. M.F and U1 have higher p-values, indicating weaker significance, but they are still included in the final model.

The R-squared value is 0.7888, meaning the model explains 78.88% of the variance in Crime rate. Adjusted R-squared is 0.7444, which considers the number of variables in the model and suggests an excellent overall fit. The F-statistic is 17.74, with a very small p-value, indicating the overall model is statistically significant.

The stepwise regression selected a subset of variables significantly impacting crime rate prediction. The final model has a high R-squared value and reduced AIC, indicating a more optimal model than the initial full-variable one.





Plot 1 checks the linearity and equal variance of the residuals. The residuals should be randomly scattered around the horizontal line at zero, and the outcomes fit.

Plot 2 checks if the residuals follow a normal distribution. The points lie along the diagonal line, indicating the residuals are normally distributed.

Plot 3 checks the constant variance of the residuals. The red line should ideally be flat. Overall, the variance of residuals in this model doesn't change across fitted values dramatically.

Plot 4 identifies observations that have a significant influence on the fitted model. Points with Cook's distance values above 0.5 or 1 (dashed lines) indicate influential observations. Thus, observations 11, 19, and 29 are the most influential in this model.

Plot 5 identifies influential points based on leverage and residual size. Observations in the top right or bottom right are points with high leverage and can significantly influence the model; thus, points 11, 19, and 29 should be reviewed closely for their impact.

Plot 6 highlights observations that are both outliers and influential. Observations outside the dashed lines indicate high influence; thus, points 11, 19, and 29 are again highlighted as significant.

Based on the plots, observations 11, 19, and 29 appear to be particularly influential. It may be worth investigating these specific data points to see if they should be removed or transformed or if the model should be adjusted to better account for them.

Method 2. Lasso

First, to make the values of the feature variables comparable, we normalized the feature variables to zero mean and unit variance to avoid the uneven influence of features of different dimensions on the model. Meanwhile, we used cross-validation to select the best lambda value and drew a graph of cross-validation error changes with lambda (figure 1) to help visualize the lambda corresponding to the smallest error. We find the regularization parameter (11.12694331514) that minimizes the model error to predict the data. The output based on the lasso model: 16 x 1 sparse Matrix of class "dgCMatrix"

```
s0
(Intercept) 905.0851064
M           84.6490497
So          22.1160730
Ed          124.1054500
Po1         308.9322037
Po2         .
LF           0.8365183
M.F         52.2295659
Pop         .
NW           4.7549442
U1          -22.1731429
U2          55.6686025
Wealth      .
Ineq        180.8039382
Prob        -81.8184334
Time        .
```

The data shows that some coefficients are 0, that is, these features do not contribute to the model (Lasso models are characterized by feature selection). $MSE = 34608.4$ represents the average squared variance between the model's prediction and the actual value. $R^2 = 0.763$ indicates that the model can explain 76.36% of the target variable fluctuations.

Finally, we plotted a scatter plot between the actual and predicted values (figure 2) to visualize the prediction effect. The red reference line indicates $y=x$, which is the ideal case where the predicted value is exactly equal to the actual value. If most of the data points are close to this line, the model prediction is good.




Figure 1 The best lambda value and the corresponding cross-validation error




Figure 2 Actual vs Predicted Crime Rates

Code:

```
library(glmnet)

crime_data <- read.table("C:/Users/Susie/Desktop/uscrime.txt", header = TRUE)

head(crime_data)

# Extract the feature variables (remove the last column)
x <- as.matrix(crime_data[, -ncol(crime_data)]) # 'Crime' is the target variable
# Extract the target variable
y <- crime_data$Crime

# Standardize the feature variables
x_scaled <- scale(x)

# Lasso regression (alpha=1 indicates Lasso regression)
lasso_model <- glmnet(x_scaled, y, alpha = 1)

print(lasso_model)

# Use cross-validation to select the optimal lambda value, determining the best regularization parameter
cv_lasso <- cv.glmnet(x_scaled, y, alpha = 1)

# Get the best lambda value (the optimal lambda selected by cross-validation)
best_lambda <- cv_lasso$lambda.min

# Print the best lambda value
print(paste("Best lambda from CV: ", best_lambda))

# Plot the cross-validation error as a function of lambda
plot(cv_lasso)

# Train the final Lasso model using the best lambda value
```

```
final_lasso_model <- glmnet(x_scaled, y, alpha = 1, lambda = best_lambda)
```

```
# Print the regression coefficients
```

```
print("Final Lasso Model Coefficients:")
```

```
print(coef(final_lasso_model))
```

```
# Make predictions using the model
```

```
predictions <- predict(final_lasso_model, s = best_lambda, newx = x_scaled)
```

```
# Calculate the Mean Squared Error
```

```
mse <- mean((predictions - y)^2)
```

```
print(paste("Mean Squared Error (MSE): ", mse))
```

```
# Define the R2
```

```
r_squared <- function(actuals, predictions) {
```

```
  ss_res <- sum((actuals - predictions)^2) # Residual sum of squares
```

```
  ss_tot <- sum((actuals - mean(actuals))^2) # Total sum of squares
```

```
  r2 <- 1 - (ss_res / ss_tot) # Calculate R2
```

```
  return(r2)
```

```
}
```

```
# Calculate R2
```

```
r2_value <- r_squared(y, predictions)
```

```
print(paste("R2: ", r2_value))
```

```
# Visualize the relationship between predicted and actual values
```

```
plot(y, predictions, main="Actual vs Predicted Crime Rates",
```

```
      xlab="Actual Crime Rates", ylab="Predicted Crime Rates", pch=19, col="blue")
```

```
abline(0, 1, col="red") # Add a reference line y=x
```

Method 3. Elastic net

We use the crime data set to build a regression model with elastic net. Before we build the model, we scale the predictors and convert it into matrix for data preparation. Then we use cross-validation to extract the best alpha value and the best lambda, based on the cross-validation error. In this case, alpha control the Lasso regularization and Ridge regularization, the lambda is regularization parameter to control the penalty.

The figure 1 shows the range of alpha with different C.V. error.

The figure 2 shows the range of lambda with different C.V. error.

The best alpha is 0.7, which means the model uses 70% Lasso and 30% Ridge.

Also, the best lambda is 6.88084, which means the coefficients will be shrunk towards zero according to both Lasso and Ridge penalties. This combination helps reduce model complexity, prevent overfitting

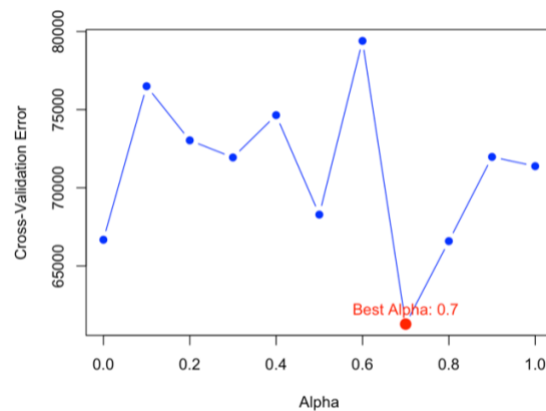


Figure 1 Cross-Validation Error vs Alpha




Figure 2 Cross-Validation Error vs Lambda (Alpha=0.7)

The best coefficients are as follow:

16 x 1 sparse Matrix of class "dgCMatrix"

s1

(Intercept) 905.085106

<i>M</i>	97.928827
<i>So</i>	19.117967
<i>Ed</i>	160.177692
<i>Po1</i>	291.352395
<i>Po2</i>	.
<i>LF</i>	.
<i>M.F</i>	57.203184
<i>Pop</i>	-9.181431
<i>NW</i>	15.074001
<i>U1</i>	-61.931326
<i>U2</i>	102.400419
<i>Wealth</i>	39.088210
<i>Ineq</i>	222.964742
<i>Prob</i>	-88.071994
<i>Time</i>	.

The data shows that some coefficients are 0, that is, these features do not contribute to the model. The R-squared is 0.7882172, which means the model can explain 78.82% of the target variable fluctuations. The figure 3 show the plot between the actual and predicted values to visualize the prediction effect. So, the model prediction is good.




Figure 3 Actual vs Predicted Crime Rates

CODE:

```
# Load necessary libraries

library(glmnet)
library(readr)

# Load the data
crime_data <- read.table("http://www.statsci.org/data/general/uscrime.txt", header = TRUE)

# Prepare predictors (X) and response (y)
x <- as.matrix(scale(crime_data[, 1:ncol(crime_data) - 1])) # Scaling the predictors and converting to matrix
y <- crime_data$Crime # Assuming 'Crime' is the dependent variable

# Cross-validation for Elastic Net model with varying alpha
cv_model <- list()
alphas <- seq(0, 1, by = 0.1)

# Loop over alpha values and store cross-validation results
cv_errors <- numeric(length(alphas))
for (i in seq_along(alphas)) {
  alpha_val <- alphas[i]
  cv_model[[as.character(alpha_val)]] <- cv.glmnet(x, y, alpha = alpha_val)
  cv_errors[i] <- min(cv_model[[as.character(alpha_val)]]$cvm) # Store the minimum cross-validation error for
  each alpha
}

# Find the best alpha based on cross-validation error
best_alpha_index <- which.min(cv_errors)
best_alpha <- alphas[best_alpha_index]
cat("Best alpha:", best_alpha, "\n")

# Plot cross-validation error vs. alpha
plot(alphas, cv_errors, type = "b", pch = 19, col = "blue", xlab = "Alpha", ylab = "Cross-Validation Error",
     main = "Cross-Validation Error vs. Alpha")
points(best_alpha, cv_errors[best_alpha_index], col = "red", pch = 19, cex = 1.5)
```

```
text(best_alpha, cv_errors[best_alpha_index], labels = paste("Best Alpha:", best_alpha), pos = 3, col = "red")
```

```
# Fit Elastic Net model using the best alpha and perform cross-validation to find the best lambda
```

```
cv_model_best_alpha <- cv_model[[as.character(best_alpha)]]
```

```
# Optimal lambda for the best alpha
```

```
best_lambda <- cv_model_best_alpha$lambda.min
```

```
cat("Best lambda:", best_lambda, "\n")
```

```
# Plot cross-validation error vs. lambda for the best alpha
```

```
plot(cv_model_best_alpha$lambda, cv_model_best_alpha$cvm, type = "b", log = "x", pch = 19, col = "blue",
```

```
      xlab = "Lambda (log scale)", ylab = "Cross-Validation Error",
```

```
      main = paste("Cross-Validation Error vs. Lambda (Alpha =", best_alpha, ")"))
```

```
points(best_lambda, min(cv_model_best_alpha$cvm), col = "red", pch = 19, cex = 1.5)
```

```
text(best_lambda, min(cv_model_best_alpha$cvm), labels = paste("Best Lambda:", round(best_lambda, 4)), pos = 3, col = "red")
```

```
# Fit the final Elastic Net model using the best alpha and lambda
```

```
elastic_net_model <- glmnet(x, y, alpha = best_alpha, family = "gaussian")
```

```
# Predict the response for the training data
```

```
y_pred <- predict(cv_model_best_alpha, s = "lambda.min", newx = x)
```

```
# Calculate  $R^2$ 
```

```
ss_res <- sum((y - y_pred)^2) # Residual sum of squares
```

```
ss_tot <- sum((y - mean(y))^2) # Total sum of squares
```

```
r_squared <- 1 - (ss_res / ss_tot)
```

```
cat("R-squared:", r_squared, "\n")
```

```
# Plot actual vs predicted values
```

```
plot(y, y_pred, xlab = "Actual Crime Rates", ylab = "Predicted Crime Rates",
```

```
      main = "Actual vs Predicted Crime Rates", pch = 19, col = "blue")
```

```
abline(0, 1, col = "red", lwd = 2) # Add 45-degree line
```



```
# Extract and print the coefficients at the best lambda  
best_coefficients <- coef(cv_model_best_alpha, s = "lambda.min")  
cat("Best Coefficients at Best Lambda:\n")  
print(best_coefficients)
```